



Statistical error ratio per bin, all bins — 2 = data_phi4seg24degFA_phifullbin (4x counts) over 1 = data_phifull (1660 shared bins)
below the dotted line the second configuration is the more precise; why each bin sits where it does: README, 'Why the statistical error differs'



Scaled as if it had more counts: data_phi4seg24degFA_phifullbin x4 (δ_{stat} by $1/\sqrt{F}$, N_{acc} by F , systematics unchanged). At nominal luminosity: data_phifull.
 δ_{stat} and systabs both carry the same $1/(f_n \cdot 0.6 \cdot 0.86)$ scaling: kinematic-dependent dilution f_n , target polarisation $P_{He} = 0.6$, effective neutron polarisation $P_n = 0.86$.
 $|A_{UT}| \cdot \text{systrel}$ does not — systrel is the quadrature of the relative uncertainties alone (3% target pol, 5% nuclear, 2.5% radiative, 3% diffractive meson, 0.2% random coincidence).